

Preregistration-first autonomous research: a protocol and a self-audit of an AI agent’s first 19 studies

[Author(s) to be decided by the lab owner]
University of Auckland

Prepared by an AI agent (Claude, Anthropic); see the AI contribution statement.

Revised draft, 6 October 2026

Abstract

AI agents can now choose a research question, analyse data and write up a result. Evidence on how such research goes wrong is accumulating as fast: post-hoc selection, overclaimed significance, underpowered nulls presented as findings, AI reviewers that accept work they have flagged, and wide variation between agent analysts.

Protocol. We describe a protocol that borrows the safeguards of confirmatory science for an autonomous research agent. One agent:

- scouts and scores research leads;
- commits each plan to a public repository before analysing outcome data;
- validates its methods;
- obtains an adversarial review from a separate model instance before any verdict;
- logs deviations;
- publishes reports that include null results.

Audit. We audit its first 19 studies (17 reported, 2 in progress) in 14 fields of the lab’s lead database. The audit uses git history, GitHub push records, coder agents of two models, and direct verification of a random sample of reported numbers.

Findings.

- Plans demonstrably preceded outcomes in 18 of 19 studies. In the first study the plan entered the public record together with its results.
- Of 16 decided verdicts, 3 (19% (7–43%)) support the preregistered hypothesis.
- 135 deviations are logged publicly, and 17 could have changed a verdict after the outcome was seen; 8 of those come from the first study.
- By the agent’s own records, which include no pre-review drafts, review changed 4 headline estimates and 2 verdicts.
- 117 of 119 sampled report numbers were verified by both coders.

Weaknesses. The audit also finds clear weaknesses:

- only one plan stated a power calculation;
- only 3 studies kept outcomes out of the agent’s reach;
- the protocol’s own rules (confirmation reviews, recorded reviewer models) were applied inconsistently;
- automated matching of report numbers to result files cannot establish that they trace.

All coders and reviewers were subagents in the executing agent’s own session. We release everything and propose changes.

1 Introduction

AI agents can now carry out every step of an empirical study: choosing a question, finding data, writing and running analysis code, interpreting results and drafting a paper [1–6].

These systems raise an old problem in a new form. Confirmatory science fixes a question and an analysis before seeing the answer, because flexible choices made afterwards can manufacture findings [7–9]. An agent can explore thousands of specifications at almost no cost, and nothing obliges it to disclose them.

Recent work documents concrete failures:

- AI scientist systems select benchmarks and runs post hoc, and audits miss this without full logs [10];
- frontier agents present underpowered negative results as findings [11];
- the interpretive statements in an AI scientist’s reports are much less accurate than its data-analysis statements [3];
- AI analysts given the same data and question reach opposite conclusions [12–14], much as human teams do [15];
- AI reviewers can be steered into accepting unsound work [16, 17].

A recent survey finds that most systems release code but few release the traces needed to verify a claim [18].

This paper. We describe and audit a research lab run by one AI agent under a protocol borrowed from confirmatory science. The agent writes each preregistration itself, before analysing outcome data, and commits it to a public git repository. The protocol was used for 19 studies between September and October 2026, all on open data, and every plan, line of code, deviation log and report is public. Our contributions are:

1. a description of the protocol as it was practised, from lead generation to publication;
2. a self-audit of whether its commitments were kept and what its safeguards caught, using git and GitHub push records, coder agents of two models, and direct verification of sampled numbers;
3. the failures the audit found, reported in the main text, with changes that follow from them.

Scope. We do not claim that the studies’ conclusions are correct; no domain expert has checked them. The audit measures process. It is also a self-audit: the agent that ran the studies designed it and wrote this paper, and every coder and reviewer was a subagent in that agent’s own session. Section 6 discusses what this does and does not allow us to conclude.

2 Related work

Autonomous research systems. Several systems now cover many fields or the whole research cycle:

- Kosmos runs hundreds of agent rollouts around a shared world model. Independent scientists judged 79.4% of its report statements accurate, but only 57.9% of its interpretive ones [3].
- Denario produced papers in more than ten disciplines [4].

- Agon runs an “omnidisciplinary” loop and contributes a taxonomy of its own failures [5].
- data-to-paper links every number in a generated paper to the code that produced it [6].
- CodeScientist and AutoDiscovery search open-endedly for findings [19, 20].
- The AI Scientist and ScientistTwo automate machine-learning research [1, 2, 21].
- Biomni executes biomedical workflows [22].
- AutoResearchClaw, an open-source system, adds citation verification [23].

The closest to preregistration is PaperClaw. It fixes a “pre-registered main-result contract” (dataset, comparison, metric and target) before experiments, and draws its verdicts from a fixed vocabulary [24]. The contract is kept in a local file rather than a public, timestamped record, and the loop halts “once the evidence supports the idea”. That stopping rule leans towards confirmation, whereas our decision rules end every study in a verdict, including null ones.

Evidence on failure. We summarised the main empirical findings above [10–14, 16, 17]. Two more strands bear on a lab like ours.

- LLM-generated ideas look more novel than human ideas but lose that advantage once executed [25], which matters for any system that ranks its own leads.
- Open health surveys have been mined for formulaic single-factor association papers without false-discovery control [26].

Critiques argue that current agent designs are not suited to autonomous discovery at all [27].

Preregistration, audit and provenance for AI research. Preregistration has been proposed for experiments *on* AI agents [28], including against LLMs not yet released [29]. Several protocols make AI-assisted research auditable.

- **Provenance and claim tracking.**
 - XScientist uses a git-like protocol with claim-to-evidence anchors and content hashes [30];
 - the Hypothesis Evolution Protocol keeps a registry of hypotheses with their lifecycles [31];
 - Atelier records every consequential exchange in files [32];
 - ResearchLoop gates each step on evidence [33];
 - Zhou and Yu combine pre-registered process observation with git sealing [34].
- **Audit of claims and discoveries.**
 - CrossAudit records audits as git commits and uses reviewers from a different vendor [35];
 - an independence-graded audit protocol ranks how independent each auditor is [36];
 - the Discovery Certification Protocol tests whether a claimed discovery can be recovered without the search history [37];
 - a post-search hold-out loop does the same for empirical economics [38];
 - ABE-Ralph checks that experiments faithfully test their claims [39];
 - claim-level auditability ties statements to their evidence [40].

Each of these was demonstrated on one or a few cases or on seeded defects. Case studies of open-source research frameworks found that none completed a full research cycle [41].

Human baseline. The human preregistration literature is the natural comparator for our deviation counts:

- in the first preregistered *Psychological Science* papers, 2 of 27 studies had no deviations and 9 disclosed none of theirs [42];
- across 459 preregistrations, 57% of publications added hypotheses and 52% omitted some [43].

What is new here. We know of no published system that does all three of the following, and no audit of such a portfolio:

- commits each study’s plan to a public, timestamped record before analysing outcome data;
- publishes a log of every deviation;
- reports the full distribution of verdicts, null results included, across a portfolio of fields.

3 The protocol as practised

The lab runs on one workstation (32-core Xeon, 94 GB RAM, one A100 80 GB GPU). The executing agent works in one long-running Claude Code session, and its model changed over time:

- from 24 September 2026, Claude Opus 5.5 (`claude-opus-5-5`). It ran studies 2–19, finished the first study, and is named in every study commit’s trailer.
- before that, for the first study’s early work (1–12 September), Claude Fable 5, Sonnet 5 and Fable 5.1. The first study’s workflow stages ran on Sonnet 5 (scoping and the plan) and Fable 5.1 (build, execution and review).

Everything it produces is committed to a public git repository and mirrored to a web page.

The human role. The lab owner sets the programme and authorises resource use, for example stopping a shared GPU service. They also make decisions the agent flags as theirs, for example whether to proceed without Māori and Pacific input on a study of an ethnicity-based waitlist policy. They wrote no plan, code or report. The first plan’s header names the owner as author, but the session logs show an agent wrote it: the executor workflow’s scope subagent, at 09:39 on 2 September 2026 (NZ time). The model logged for it is Claude Sonnet 5.

The first study was different. It ran under an executor workflow (`run-avenue.js`) with two reviewers, a methodology lens and a code lens, followed by a third review and a final audit. Studies 2 onwards ran in the main session with one reviewer. We describe the protocol as it settled after the first study, and report the first study separately where it matters.

3.1 Leads: scouting and triage

Leads come from three scouting rounds run as multi-agent workflows:

- 10 domain scouts, a completeness critic and 2 gap-fill scouts;

- 8 scouts and a calibration critic;
- 4 local scouts on Auckland and Aotearoa New Zealand, with a critic.

They are stored in `leads.json`. Each lead records:

- a hypothesis;
- the gap it fills;
- literature for and against;
- open datasets;
- a method and plan;
- risks;
- value and cost scores (1–10);
- a critic’s note;
- a plain-language summary.

There are 92 leads in 21 fields. The Auckland leads count as one field (`auckland-nz`), although Table 1 gives their subject areas.

The agent works through the high-value, low-cost quadrant of a published value-against-cost plot. Before a lead is executed it gets a fresh novelty check. One lead (`software-security-attestation-phantom-code`) was dropped this way because a 2026 paper had pre-empted it.

3.2 Preregistration

Before analysing outcome data, the agent writes `plan.md`, then commits and pushes it. The plan contains:

- the question and hypothesis;
- a disclosure of what has already been seen;
- data and units;
- the primary statistic;
- a decision rule mapping results to a verdict;
- validation steps;
- secondary analyses, labelled as such.

The verdicts are *Supported*, *Refuted*, *Unsupported*, *Inconclusive* and *Pending* (for prospective tests). The rule must say when the evidence is too weak to decide.

A commit shows when a plan was *committed*, not when it was *written* or whether the data had been explored first. That gap is narrowest when outcomes are kept out of reach.

3.3 Keeping outcomes out of reach, and verification

Where a design allows, outcomes are kept from the agent until the analysis is fixed.

- **Sealed outcomes.** The data are stored unread with only a SHA-256 hash committed. Examples are the post-period crash data in the speed-limit study and the quantum-hardware answer strings in the peaked-circuit study. The sealed files stay local; releasing them after the study is one of our proposed changes.
- **Prospective outcomes.** The outcome does not yet exist, for example City Rail Link ridership through March 2027.

Otherwise the plan’s disclosure is the only safeguard.

Separately from blinding, one study’s outcomes are machine-verified: the formal-mathematics study’s vacuity certificates are checked by the Lean kernel.

3.4 Validation, review and confirmation

Validation. Methods are validated before use on synthetic data with known truth, positive and negative controls, or a reproduced published number.

Independent review. Before a verdict, a separate agent instance spawned in the same session reviews the plan, code, results and draft report, and is told to try to refute them. The reviewer’s model is recorded in the deviation logs of 10 of the 17 reported studies, as Claude Fable 5.1 (`claude-fable-5-1`); for the others the record does not say. The first study’s reviewers returned sound/fixable/refuted; later reviewers returned “fix first” or “publish with edits”.

Confirmation. The intended practice is that the same reviewer checks the fixes after a “fix first”. This happened in 6 of the 12 studies whose review demanded fixes.

Pre-review drafts. No study committed its pre-review draft. What review changed is therefore known only from the agent’s own records (Section 5.5).

3.5 Deviation log and reporting

Deviation logs. `deviations.md` records departures from the plan, normally with a timestamp from the system clock. In the first study, 26 of its 46 section headers carry no timestamp of their own.

Reports. Each report opens with an “In plain terms” section; all 17 are within the target of 150–300 words (224–299). Only 7 reports have a section summarising the review; the others mention it in a preamble and in passing. The verdict, a note and a headline are written back to `leads.json` and the public page.

4 Self-audit design

The audit protocol (`paper/audit_plan.md`) was committed at `f424e29`, before any measure was computed. The audit code and the first automated outputs followed six minutes later (`ea1736b`), so the code was written alongside the protocol, not committed with it as the protocol states. Every departure from the protocol is logged in `paper/audit/deviations.md` (A1–A10), including those made in response to the review.

Coders. Judgement-based coding used two coder agents, both subagents in the executing agent’s session:

- coder A: Claude Fable 5.1;
- coder B: Claude Opus 5.5, the executing agent’s own model.

They saw only the study files and a fixed codebook, and not each other’s work. Their prompts are published verbatim (`paper/audit/coder_prompts.md`). A third coder (Opus 5.5) made an inventory of the safeguards in each plan.

Measures.

- (A) Whether each plan was committed and pushed before the first outcome file. This uses git and the GitHub push events, archived in `audit/`.
- (B) Every deviation entry in the reported studies, coded by category, trigger, timing and verdict relevance.
- (C) What independent review raised and changed.
- (D) Whether the numbers in reports are accurate. This uses direct verification of a random sample, with automated matching as a secondary measure.
- (E) Whether cited arXiv identifiers and DOIs resolve.
- (F) Scout estimates against outcomes.

Agreement is reported as Cohen’s κ [44], and proportions with Wilson 95% intervals.

5 Results

5.1 The portfolio

Table 1 lists the 19 studies in the order their plans were committed. They cover 14 of the 21 fields in the lead database.

The first study ran from early September to 25 September 2026 under the original executor workflow. The other 16 reports appeared between 1 and 4 October, with several studies running in parallel.

Of the 16 decided verdicts:

- 3 support the preregistered hypothesis (19% (7–43%));
- 3 refute it;
- 2 are *Unsupported*: the evidence points against the hypothesis without meeting the refutation rule;
- 8 are *Inconclusive* (50% (28–72%)).

One prospective test is *Pending* until its outcome data exist in 2027.

5.2 Were plans committed before outcomes?

Commit order. Git shows the plan committed strictly before any results file or report in 15 of 19 studies. Across those 15, the median gap to the first output was 3.5 hours (range 0.45–20.4 h); 3 plans came less than an hour before their first output.

Push records. GitHub’s server-side push events, archived in `audit/`, back up the commit timestamps:

- 17 of the 19 plan commits were pushed within a minute of their commit time;
- one (bioinformatics) was pushed after 22 minutes;
- the first study’s commits were pushed 5.6 days after they were made;
- all 15 plans committed strictly first were also pushed before their first output.

The four same-commit cases. In the other 4 studies the plan and some results files arrived in a single commit. Both coders classified every file in those commits.

- **Three contained only pre-outcome material:** pre-period tables and the hash of a sealed post-period file (speed limits), a catchment inventory (wastewater), and a data-coverage table (FluSight).
- **The fourth was the first study (GWTC-4).** Its plan entered git in the same commit as the full results, including the verdict.

The first study’s plan file did exist earlier: the agent’s private session record shows it, with the committed content, on 2 September 2026. That record is not public or externally timestamped, so it is weaker evidence than the commit history. The report of that study described itself as preregistered, and it now carries a correction.

Overall. Counting pre-outcome material as compliant, plans demonstrably came first in 18 of 19 studies (95% (75–99%)). No plan was edited after its first commit. This shows only that plans were *committed* first. It cannot show that they were written before the data were explored, which is what the disclosure sections and, where used, sealing are for.

5.3 Which safeguards did the plans use?

- **Applied consistently.** Every plan stated a decision rule. 18 disclosed what had already been seen (and 1 partially; an earlier text search found explicit sections in only 13, because it looked for specific headings). 18 used only data that are open without registration.
- **Applied unevenly.** 13 plans required validation before the primary analysis, 4 did so partially and 2 did not. Of the 8 plans with several primary tests, 2 addressed multiplicity. These codes come from a single coder and carry no reliability estimate.
- **Power.** Only one plan (speed limits) stated a minimum detectable effect before results, and one more did so partially. Two studies record both an assumed effect and a detectable one:
 - the speed-limit study hypothesised a crash rate ratio of 1.10 but could detect only 1.33;
 - the wastewater study assumed a ratio of 0.85 and, after review, computed a detectable ratio of 0.71.

Both studies were underpowered for the effect they set out to detect, and both ended *Inconclusive*.

- **Keeping outcomes out of reach.** This was done in only 3 studies: two sealed (speed limits, quantum) and one prospective (City Rail Link). 2 more did so partially. In the other 14, the outcome data were public and the disclosure section is the only safeguard. Only one *decided* study had its outcomes sealed.

5.4 Deviations

The coders identified 141 and 142 deviation entries across all 19 studies. The difference is whether one review paragraph in GWTC-4 counts as a separate entry, and we kept coder A’s segmentation. Following the audit plan, the analysis covers the 135 entries in the 17 reported studies.

Agreement. Agreement was high:

- category $\kappa = 0.87$;
- trigger $\kappa = 0.90$;
- timing $\kappa = 0.96$;
- verdict relevance $\kappa = 0.83$.

Disagreements were settled with three fixed rules: category follows the cause; verdict relevance takes the cautious side; timing is “unclear” if either coder said so. Each decision is logged.

Categories and triggers. The categories were:

- added secondary analyses: 31;
- clarifications: 30;
- bug or numerical fixes: 22;
- environment and logistics: 15;
- changes to the primary specification: 13;
- data problems: 14;
- withdrawn or corrected claims: 8;
- decision-rule changes: 2 (Figure 1).

73 entries were made after the primary outcome was seen. 38 of these were requested by the reviewer, which necessarily comes after the results, and 35 were initiated by the agent or by a validation check.

Verdict relevance. 29 entries (21% (15–29%)) could plausibly have changed the verdict, and 17 of them were made after the outcome was seen. Of those:

- 11 were triggered by validation or review, mostly corrections of errors;
- 6 were the agent’s own changes after seeing the outcome.

The first study accounts for 36 of the 135 entries and 8 of the 17 verdict-relevant post-outcome ones. It also accounts for most of the agent’s own post-outcome changes (entries D19, D20, D21, D29): a primary statistic redefined, a statistic abandoned, a decision rule chosen after the real-data values were seen, and a recalibration. Without the first study:

- 99 entries remain;
- 16 are verdict-relevant (16% (10–25%));
- 9 are verdict-relevant and post-outcome;
- 2 of those were the agent’s own changes: a corrected claim and a numerical guard.




Figure 1: Logged deviations in the 17 reported studies, by category and trigger (adjudicated codes).

Decision rules. Two decision rules were changed:

- GWTC-4’s, after the outcome was seen; the study’s own final audit judged it post hoc;
- the City Rail Link study’s, after review but before any outcome data existed. A further City Rail Link entry moved a threshold and was coded as a specification change.

Comparison with human studies. In human preregistered studies, undisclosed deviations are common: 9 of 27 early *Psychological Science* studies disclosed none of theirs [42]. Here every deviation the agent made is, by construction, in a public log. The audit cannot detect deviations that were never logged.

5.5 What independent review changed

Recommendations. Every review asked for changes:

- 11 said “fix first”;
- 1 (the first study’s methodology review) said “refuted”;
- 5 said “publish with edits”.

A confirmation pass followed in 6 of the 12 studies whose review demanded fixes.

Issues raised. The coders agreed on the exact number of issues per review only 24% of the time, so we give both totals: 127 and 135. By category:

- analysis or model errors: 26 and 34;
- computational errors: 11 and 18;
- overclaiming: 31 and 29;
- missing analyses: 25 and 30;
- documentation: 32 and 22;
- leakage or blinding: 2 each.

What changed. No study committed a pre-review draft, so what review changed is known only from the agent’s own deviation logs and the reports’ accounts of their reviews. On that evidence:

- **Headlines.** Review changed the headline estimate in 5 studies according to coder A, 5 according to coder B, and 6 according to either coder; we adjudicated 4.
- **Verdicts.** Review changed the verdict in 3 (A), 2 (B) or 3 (either); we adjudicated 2.

These measure-C adjudications followed no rule set in advance. In each case the adjudicated count is the lower one, and we report all three counts for that reason.

The clearest cases, each documented in the run directory, were these:

- **Seismology:** the reviewer found a local optimum, and the Norcia recovery fraction moved from 0.51 to 0.88. The pre-review value is in `results/tables/primary_round1.csv`.
- **GWTC-4:** the verdict went from “confirmed”, recorded in the commit message of the plan commit, to “indeterminate”. Its headline false-positive rate was replaced: a different statistic, not a re-estimate.
- **Citation-context study:** the verdict went from a refutation to “uninformative” after its positive control failed.

5.6 Are the numbers in reports accurate?

Direct verification. Both coders independently checked a seeded random sample of 119 substantive numbers, 7 from each report, locating the specific source of each.

- **Consistent with their source:** 119 (coder A) and 117 (coder B); 117 by both coders (98% (94–100%); agreement on the class $\kappa = 0.95$). Sources ranged from a results cell to a recomputation from the data or a design parameter.
- **Inconsistent:** 2. Both were flagged by coder B; coder A had judged them to sit on a rounding boundary. Recomputation confirmed both as rounding errors of one unit in the last printed digit, caused by rounding twice: 0.9149 printed as 0.92, and 0.35497 printed as 0.36. The first also appears in a deviation entry.
- **Not verifiable:** 0 could not be verified by at least one coder.

A coder also found a phrase in one report that described a share of boardings as if it were a share of routes. Every confirmed error has been corrected in the report and logged in the study’s deviation file.

Untraced sample. An earlier sample of 85 numbers that automated matching had failed to trace gave a similar picture: no number inconsistent with the results, and three untraceable. All three untraceable numbers are figures the reviewer computed in scratch code that was never saved. Outside that sample, one coder found a third rounding error by the same mechanism, now corrected. The sample is lopsided: 3 reports contributed 10 numbers each, out of 15 reports in all.

Automated matching does not work. We first tried to trace report numbers automatically. The method searches a study’s `results/` files for any value within rounding of each printed number.

- It matched 92% of substantive numbers. Shifting each number by 3–7 units in its last printed digit and searching the same files still matched 75% (Figure 2).
- At three significant figures the rates were 92% and 67%.




Figure 2: Share of report numbers that match some value in the study’s **results/** files, for the printed value (blue) and for the same value shifted by 3–7 units in its last digit (gold). The size-matched chance rate is close to the rate for the real numbers, so a match carries little information.

Result files contain so many nearby values that a match carries little information, and a wrong number can match by chance. Earlier controls that matched against other studies’ files understated this. Tracing a claim needs a pointer to a specific cell, not a search for a matching value.

What the automated pass did reveal is 19 untraced numbers with five or more significant figures. 16 of them are sample sizes from one study (bioinformatics) that appear nowhere in its results files.

5.7 References

All 169 arXiv identifiers cited in leads, plans and reports resolve. 109 of 111 DOIs resolve automatically, and the other two do once parsing errors are repaired. The 33 arXiv identifiers and 6 DOIs in this paper’s own bibliography were checked the same way. Resolution shows that a reference exists, not that it says what it is cited for.

5.8 Scout estimates against outcomes

Effort. The scouts estimated a median of 5.0 days of wall time per study (84 days in total).

- **The first study** took far longer than estimated. Its transcript activity spans 34 days against an estimate of 3–6, with interruptions.
- **The later studies** took far less. The 16 later reports appeared within 3.4 days, with several running in parallel.

Per-study effort cannot be measured reliably, because commits were made in bursts and transcripts rarely name the study they concern. A transcript-based estimate gives 38 active hours across all studies, and we do not rely on it.

Lead scores. Every executed lead had a value score of 6 or 7 and a cost score between 2 and 5. With so little variation, nothing can yet be said about whether the scores predict verdicts.

6 Discussion

What the protocol delivered. The central commitment held for the 18 studies after the first: each plan was committed and pushed publicly before outcome files appeared, and none was edited afterwards.

Null and inconclusive verdicts were reported as such. Only 3 of 16 decided verdicts (19% (7–43%)) support the preregistered hypothesis. For context, about 44% of registered reports in psychology support theirs [45]. Our interval reaches that value, the categories differ, and our hypotheses were written by the same agent that tested them, so the comparison is only suggestive. Even so, the distribution is the opposite of the curated lists of discoveries by which autonomous systems are usually presented.

The decision rules alone did not produce this:

- in the record-margin study, an unplanned validation exposed a biased null model and turned a preregistered “supported” into “inconclusive”;
- in the bioinformatics study, review restored unfavourable results that the draft had omitted.

Validation and review did that work.

Review was not a formality. Every review demanded changes, and the deviation logs record 135 departures from plan. For comparison, most early human preregistrations disclosed few or none of their deviations [42].

What did not work.

- **Committed first is not written first.** In 14 of 19 studies the outcome data were public and within the agent’s reach before the plan was committed, sometimes by under an hour. The only protection is the agent’s own disclosure of what it had seen. Sealing outcomes was the exception (3 studies), not the rule.
- **Missing power calculations.** Only one plan stated a minimum detectable effect. The two studies that recorded both an assumed and a detectable effect were underpowered by a wide margin, and half the decided verdicts are *Inconclusive*. Kirgis et al. report the same weakness in frontier agents [11].
- **Rules applied unevenly.** The protocol’s own rules were not applied consistently:
 - confirmation passes followed 6 of 12 demands for fixes;
 - the reviewer’s model is recorded for 10 of 17 studies;
 - only 7 reports summarise their review;
 - no study committed a pre-review draft, so the effect of review rests on the agent’s own account.
- **The first study.** It was run under an earlier, looser process. It committed its plan with its results (although the plan existed three weeks earlier), recorded a “confirmed” verdict before review, and made the portfolio’s most consequential post-outcome changes, including a decision rule chosen after seeing the data. The protocol as described here stabilised after it.
- **Number accuracy.** Direct verification found the sampled numbers overwhelmingly consistent with their sources. It also found that one error mechanism, rounding twice, recurred

in three reports. Automated matching of values to result files, which we and others had recommended, turned out to be nearly uninformative.

Limitations.

- **This is a self-audit.** The agent that ran the studies designed the audit, wrote this paper and adjudicated disagreements. Every coder and the reviewer of this paper were subagents spawned in its own session, from prompts it wrote. One coder shares its model, and all are from one vendor. The reviewer of this paper did find errors the agent had missed, including statistics computed over the wrong set, a miscounted recommendation, and missing prior work. But this falls short of the independence that cross-vendor or human audit provides [35, 36].
- **Self-attested records.** The agent controls the repository and its own session records. GitHub’s push records corroborate commit times, but they cannot show when a plan was written.
- **No check of truth.** No domain expert checked the conclusions. The audit measures process.
- **A small, uneven sample.** Nineteen studies from one lab, one agent and one month are a case series. The first study differs from the rest in kind.

Changes to the protocol. The audit and its review point to the following changes:

1. register each plan with an external, timestamped registry as well as git, and state in it what data the agent holds;
2. require a minimum detectable effect, and an equivalence-type rule for *Refuted*, in every plan, and scout leads with power in mind;
3. seal outcomes whenever a design allows, and release the sealed files when the study ends;
4. commit the pre-review draft before requesting review, and publish the reviewer’s prompt and full review;
5. make the confirmation pass mandatory after any demand for fixes, and record every model by identifier;
6. link every number in a report to a specific result cell, and round only once, from the stored value;
7. for *Supported* verdicts, add a reviewer from another vendor or a human, plus a blind re-analysis by a second analyst agent given only the plan and the data [12, 13];
8. record scout estimates against outcomes, and state in every report which decisions a human made.

AI contribution statement

Everything in this lab was produced by AI agents (Anthropic models):

- **Leads.** Generated by scout and critic agents.
- **Studies.** Every plan, line of code, deviation log and report was written by the executing agent or its workflow subagents. That was Claude Opus 5.5 from 24 September; the first

study began under Fable 5, Sonnet 5 and Fable 5.1, and Sonnet 5 wrote its plan. Study-level reviews were by Claude Fable 5.1 subagents.

- **This paper.** The executing agent designed the audit and wrote the paper.
- **Audit coding.** Coder A (Fable 5.1), coder B (Opus 5.5) and an inventory coder (Opus 5.5) did the judgement-based coding.
- **Review of the paper.** A Fable 5.1 subagent reviewed the paper.

All coders and reviewers were spawned by the executing agent within its own session, from prompts it wrote. They are therefore not independent in the sense used by cross-vendor or human audit. The review of this paper is published unedited ([paper/review/review.md](#)), and Appendix B lists how each point was handled.

The human lab owner set the programme, approved resource use and decisions flagged to them, and takes responsibility for any submission. No human checked the conclusions of the individual studies.

Data and code availability

Public. Everything listed here is in <https://github.com/uoa-ereseach/citations> under `research-lab/`:

- all plans, code, deviation logs and reports;
- the audit protocol, code, coder prompts, coder outputs and adjudications;
- the archived GitHub push events;
- this paper’s review.

Raw data can be downloaded again from the sources each plan names.

Not public. Four things are withheld:

- the session transcripts, because they contain the lab owner’s messages;
- the coders’ transcripts;
- the study-level reviewer prompts and full reviews, which were not saved for studies 1–17;
- two sealed outcome files, which remain local.

References

- [1] Chris Lu, Cong Lu, Robert Tjarko Lange, Yutaro Yamada, et al. Towards end-to-end automation of AI research. *Nature*, 651, 2026. doi:10.1038/s41586-026-10265-5; the AI Scientist, first described in arXiv:2408.06292.
- [2] Yutaro Yamada, Robert Tjarko Lange, Cong Lu, et al. The AI scientist-v2: Workshop-level automated scientific discovery via agentic tree search. arXiv:2504.08066, 2025.
- [3] Ludovico Mitchener, Angela Yiu, Benjamin Chang, et al. Kosmos: An AI scientist for autonomous discovery. arXiv:2511.02824, 2025.

- [4] Francisco Villaseca-Navarro, Boris Bolliet, Pablo Villanueva-Domingo, et al. The Denario project: Deep knowledge AI agents for scientific discovery. *arXiv:2510.26887*, 2025.
- [5] Youran Sun, Xingyu Ren, Chugang Yi, Jiaxuan Guo, Kejia Zhang, Jianda Du, and Haizhao Yang. Agon: An autonomous large-scale omnidisciplinary research system built on prompt economy. *arXiv:2606.24177*, 2026.
- [6] Tal Ifargan, Lukas Hafner, Maor Kern, Ori Alcalay, and Roy Kishony. Autonomous LLM-driven research—from data to human-verifiable research papers. *NEJM AI*, 2025. doi:10.1056/AIoa2400555.
- [7] Joseph P. Simmons, Leif D. Nelson, and Uri Simonsohn. False-positive psychology: Undisclosed flexibility in data collection and analysis allows presenting anything as significant. *Psychological Science*, 22(11):1359–1366, 2011.
- [8] Andrew Gelman and Eric Loken. The statistical crisis in science. *American Scientist*, 102(6):460–465, 2014.
- [9] Brian A. Nosek, Charles R. Ebersole, Alexander C. DeHaven, and David T. Mellor. The preregistration revolution. *Proceedings of the National Academy of Sciences*, 115(11):2600–2606, 2018.
- [10] Ziming Luo, Atoosa Kasirzadeh, and Nihar B. Shah. The more you automate, the less you see: Hidden pitfalls of AI scientist systems. *arXiv:2509.08713*, 2025.
- [11] Peter Kirgis, Sayash Kapoor, Andrew Schwartz, et al. Can AI agents conduct open-ended AI research? Early evidence from two case studies. *arXiv:2607.27191*, 2026.
- [12] Martin Bertran, Riccardo Fogliato, and Zhiwei Steven Wu. Many AI analysts, one dataset: Navigating the agentic data science multiverse. *arXiv:2602.18710*, 2026.
- [13] Ruijiang Gao and Steven Chong Xiao. Nonstandard errors in AI agents. *arXiv:2603.16744*, 2026.
- [14] Jiacheng Miao, Jonathan K. Pritchard, and James Zou. The agentic garden of forking paths. *arXiv:2607.01507*, 2026.
- [15] Albert J. Menkveld et al. Nonstandard errors. *The Journal of Finance*, 79(3):2339–2390, 2024.
- [16] Fengqing Jiang, Yichen Feng, Yuetai Li, et al. BadScientist: Can a research agent write convincing but unsound papers that fool LLM reviewers? *arXiv:2510.18003*; *ACL 2026*, 2025.
- [17] Federico Bianchi, Owen Queen, Nitya Thakkar, et al. Exploring the use of AI authors and reviewers at Agents4Science. *arXiv:2511.15534*, 2025.
- [18] Tianyu Ding, Aditya Nannapaneni, Bingfan Liu, et al. Autonomous research agents: A survey of AI scientists and the verification gap. *arXiv:2608.05179*, 2026.
- [19] Peter Jansen, Oyvind Tafjord, Marissa Radensky, et al. CodeScientist: End-to-end semi-automated scientific discovery with code-based experimentation. Allen Institute for AI; *arXiv:2503.22708*, 2025.
- [20] Dhruv Agarwal, Bodhisattwa Prasad Majumder, Reece Adamson, et al. AutoDiscovery: Open-ended scientific discovery via Bayesian surprise. *arXiv:2507.00310*, 2025.

- [21] Jaehyun Nam, Jinsung Yoon, Yanzhou Pan, Yubo Wang, Rui Meng, Parthasarathy Ranganathan, and Tomas Pfister. ScientistTwo: Pioneering the human knowledge frontier with autonomous AI. *arXiv:2609.19644*, 2026.
- [22] Kexin Huang et al. Autonomous biomedical research with an artificial intelligence agent. *Science*, doi:10.1126/science.adz4351, 2026.
- [23] Jiaqi Liu, Shi Qiu, Mairui Li, et al. AutoResearchClaw: Self-reinforcing autonomous research with human-AI collaboration. *arXiv:2605.20025*, 2026.
- [24] Weiwei Ye, Hangchen Liu, Dongyuan Li, and Renhe Jiang. PaperClaw: Harnessing agents for autonomous research and human-in-the-loop refinement. *arXiv:2606.22610*, 2026.
- [25] Chenglei Si, Tatsunori Hashimoto, and Diyi Yang. The ideation–execution gap: Execution outcomes of LLM-generated versus human research ideas. *arXiv:2506.20803*, 2025.
- [26] Tulsi Suchak et al. Explosion of formulaic research articles, including inappropriate study designs and false discoveries, based on the NHANES US national health database. *PLOS Biology*, 23:e3003152, 2025.
- [27] Harshit Bisht, Vinay Kumar, Kevin Maik Jablonka, et al. Agentic AI scientists are not built for autonomous scientific discovery. *arXiv:2605.08956*, 2026.
- [28] Michelle Vaccaro. Preregistration for experiments with AI agents. *arXiv:2606.11217*, 2026.
- [29] Maria Thomas, Kristina Gligoric, and Nihar B. Shah. Mitigating LLM-based p-hacking by preregistering for the next LLM. *arXiv:2606.27687*, 2026.
- [30] Jixiang Luo. XScientist: A git-like research protocol for long-running autonomous scientific discovery. *arXiv:2607.12301*, 2026.
- [31] Izumi Takahara and Teruyasu Mizoguchi. Toward auditable AI scientists: A hypothesis evolution protocol for LLM agents. *arXiv:2607.09195*, 2026.
- [32] Geng and Lu. Atelier: A human-owned protocol for auditable AI-assisted research. In *International Conference on Machine Learning (ICML)*, 2026.
- [33] Yihan Xia and Taotao Wang. ResearchLoop: An evidence-gated control plane for AI-assisted research. *arXiv:2605.28282*, 2026.
- [34] Yang Zhou and Chengqun Yu. Auditable AI-assisted research writing: An engineering discipline with pre-registered process observation. *arXiv:2608.10858*, 2026.
- [35] Zhaohe Dong and Yuhao Chen. CrossAudit: A git-native, cross-vendor audit loop for agentic science. *arXiv:2608.28631*, 2026.
- [36] Mohamed Chahine Ghanem. Who audits whom, on what substrate, with what evidence? An independence-graded audit protocol for agentic AI. *arXiv:2609.18272*, 2026.
- [37] Jingjie Ning, Shanshan Zhong, Xiaochuan Li, and Ji Zeng. Scores alone do not prove discovery: The discovery certification protocol for auditing AI research agents. *arXiv:2609.09219*, 2026.
- [38] Minchul Shin. An auditable AI agent loop for empirical economics: A case study in forecast combination. *arXiv:2603.17381*, 2026.

- [39] Lezhi Yu, Xiaogang Xu, Yuhua Zhou, Shuibing He, and Aimin Pan. Beyond execution: Auditing experimental fidelity in LLM-driven scientific research. arXiv:2608.26753, 2026.
- [40] Razeen A. Rasheed, Somnath Banerjee, Animesh Mukherjee, and Rima Hazra. From fluent to verifiable: Claim-level auditability for deep research agents. arXiv:2602.13855, 2026.
- [41] Shreyansh Agrawal et al. Can AI conduct autonomous scientific research? Case studies on two real-world tasks. bioRxiv doi:10.64898/2026.01.05.697809, 2026.
- [42] Aline Claesen, Sara Gomes, Francis Tuerlinckx, and Wolf Vanpaemel. Comparing dream to reality: an assessment of adherence of the first generation of preregistered studies. *Royal Society Open Science*, 8:211037, 2021. doi:10.1098/rsos.211037.
- [43] Olmo R. van den Akker et al. Selective hypothesis reporting in psychology: Comparing preregistrations and corresponding publications. *Advances in Methods and Practices in Psychological Science*, 6(3), 2023. doi:10.1177/25152459231187988.
- [44] Jacob Cohen. A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1):37–46, 1960.
- [45] Anne M. Scheel, Mitchell R. M. J. Schijen, and Daniël Lakens. An excess of positive results: Comparing the standard psychology literature with registered reports. *Advances in Methods and Practices in Psychological Science*, 4(2), 2021.

A Where to find each part of the protocol

All paths are relative to **research-lab/** in the public repository.

Artifact	Location and content
Lead database	<code>leads.json</code> : hypothesis, gap, literature for and against, datasets, method, plan, risks, value, cost, feasibility, critic note, plain-language summary, and (after execution) status, verdict tag, verdict note and headline
Lead explorer	<code>index.html</code> : value against cost plot (GitHub Pages)
Scouting workflows	<code>research-lead-scout-round2.js</code> and earlier rounds; <code>run-avenue.js</code> (executor workflow, used in early studies)
Per-study plan	<code>runs/<id>/plan.md</code> , committed before outcome analysis
Deviation log	<code>runs/<id>/deviations.md</code> , entries timestamped with the system clock
Report	<code>runs/<id>/report.md</code> and <code>report.html</code> , with “In plain terms”, results, caveats and review summary
Code and results	<code>runs/<id>/code/</code> , <code>runs/<id>/results/</code>
Field survey	<code>related-efforts.md</code>
Self-audit	<code>paper/audit_plan.md</code> (protocol), <code>paper/audit/</code> (all measures, coder outputs, adjudication, deviations), <code>paper/code/</code> (audit and paper-number code)

B Independent review of this paper

The paper was reviewed by a Claude Fable 5.1 subagent spawned in the executing agent’s session, using the prompt in `paper/review/prompt.md`. Its review is published unedited as `paper/review`

/review.md. It recommended *major revision*, raising 9 major and 23 minor issues, and recomputed about 60 of the paper’s numbers. We list each major issue and what we did.

M1. Review effects rest on self-reports. Accepted. Section 5.5 now says that no study committed a pre-review draft, so measure C rests on the agent’s own records. It cites the evidence for the clearest cases. Committing the pre-review draft is now protocol change 4.

M2. Protocol rules not followed as described. Accepted. Section 3 now describes the protocol as practised, with counts measured from the run directories:

- confirmation passes (6 of 12);
- recorded reviewer model (10 of 17);
- review sections (7 of 17);
- headers without timestamps in the first study;
- the first study’s two-reviewer scheme;
- the change in recommendation vocabulary.

“Formal” verification is no longer listed as blinding, and the executing model is named.

M3. Undisclosed departures from the audit plan. Accepted. We made these changes:

- added Wilson intervals;
- removed the in-progress studies from measure B, as the plan specified;
- corrected the statement that the audit code was committed with the plan;
- reported both coders’ measure-C counts beside the adjudicated ones, and said that those adjudications followed no rule set in advance;
- reverted coder A’s raw file to its committed state;
- had both coders classify the four same-commit cases (they agreed with each other and with the earlier classification);
- reported the power comparison.

M4. Wrong or misleading statistics. Accepted. We corrected:

- the median plan-to-output gap, now computed over the 15 studies with the plan strictly first (3.5 h, not 1.7 h);
- the recommendation count (11 “fix first”, 1 “refuted”, 5 “publish with edits”);
- the comparison with registered reports, now given with its interval and caveats;
- the GWTC-4 false-positive rate, described as replaced rather than re-estimated;
- the effort claim, now restricted to the later studies and stating the first study’s overrun;
- the lead-score claim, now noting the restricted range.

M5. Novelty and related work. Accepted. We now cite PaperClaw and describe how its local “pre-registered contract” differs. We also cite XScientist, the Hypothesis Evolution Protocol, Atelier, Zhou and Yu, ResearchLoop, the independence-graded audit protocol, claim-level auditability and AutoResearchClaw, and the human preregistration adherence baseline. The novelty claim is narrowed. Uncited bibliography entries are now cited, and the Nature reference is corrected.

M6. Independence overstated. Accepted. The AI contribution statement and Section 4 now say the following:

- the coders and this reviewer were same-session subagents working from prompts written by the executing agent;
- which model each coder used;
- that one coder shares the executing agent’s model.

The coder prompts are published.

The working tree was frozen during the confirmation pass. The edits made during the first review are committed, and the reviewer’s “version reviewed” note records them.

M7. Number tracing weaker than presented. Accepted, and the measure redesigned. A size-matched perturbation control shows that automated value matching is nearly uninformative. The primary measure is now direct verification of a random sample of 119 numbers by both coders. That verification found real errors, which have been corrected in the reports. The untraced five-significant-figure numbers and the lopsided untraced sample are now described.

M8. Push events. Accepted, with one correction. The archived push events show that 17 plans were pushed within a minute. The bioinformatics plan was pushed after 22 minutes, and the first study’s commits after 5.6 days, so the reviewer’s “all 19 within 0.1 minutes” does not hold for the first study. The distinction between “committed first” and “written first” is now stated in Sections 3 and 5.2.

M9. The first study is different. Accepted. It is described separately in Section 3. Deviation statistics are given with and without it, and its post-outcome changes are listed. We added the evidence from the private session record that its plan existed on 2 September.

Minor issues. All 23 were addressed by editing the text, except two:

- **Minor 4** (the first plan’s author header) is resolved. The lab owner confirmed that an agent wrote the plan, and the session logs identify it as the workflow’s scope subagent, running Claude Sonnet 5.
- **Minor 23** (plain-language word counts) is reported as a check that passed.

A confirmation pass by the same reviewer is recorded in `paper/review/`.

Table 1: The 19 studies in order of plan commit. The “Field” column gives the subject area; the lead database groups the five Auckland studies as one field. “Plan first”: git shows the plan committed before any outcome file; [†]plan committed together with pre-outcome material only (both coders agree; Section 5.2). “Blinding”: how outcomes were kept from the agent. Sealed means stored unread with a hash committed; prospective means the outcome did not yet exist; partial means a pending hold-out (waitlist) or published final scores already known (MCF). “Power”: whether the plan stated a power calculation or minimum detectable effect.

#	Field	Question	Verdict	Plan first	Blinding	Power
1	Astronomy	Is the GWTC-4 mass-ratio-spin anticorrelation real?	Inconclusive	no	–	–
2	Bioinformatics	Do ClinGen variant-predictor thresholds still hold?	Refuted	yes	–	–
3	Climate	Are heat records broken by growing margins?	Inconclusive	yes	–	–
4	Health economics	Does flu wastewater improve hospitalisation forecasts?	Unsupported	same commit [†]	–	–
5	Metascience	Does citing language predict replication outcomes?	Inconclusive	yes	–	–
6	Climate	Does moisture treatment change AI humid-heat attribution?	Refuted	yes	–	–
7	Computational social science	Did Wikipedia temporary accounts raise vandalism?	Inconclusive	yes	–	–
8	Neuroimaging	Is shrinking small fMRI maps to a meta-analytic prior worth it?	Refuted	yes	–	–
9	ML science	Does answer-letter mass predict multiple-choice format emergence?	Unsupported	yes	partial	–
10	Cheminformatics	Are benchmark activity cliffs partly measurement noise?	Supported	yes	–	–
11	Transport (NZ)	Is City Rail Link rider-ship uplift structural?	Pending	yes	prospective	partial
12	Transport (NZ)	Did reversing speed-limit cuts increase crashes?	Inconclusive	same commit [†]	sealed	yes
13	Public health (NZ)	Is COVID-19 under-ascertainment greater in deprived towns?	Inconclusive	same commit [†]	–	–
14	Urban environment (NZ)	Did Auckland’s 2016 upzoning cost tree canopy?	Inconclusive	yes	–	–
15	AI evaluation	Do independent audits of SWE-bench Verified agree?	Supported	yes	–	–
16	Health policy (NZ)	Did ending an ethnicity-inclusive waitlist score change waits? 21	Inconclusive	yes	partial	–
17	Seismology	Do neural point processes beat ETAS only via incompleteness?	Supported	yes	–	–
18	Quantum simulation	Can one GPU re	in progress	yes	sealed	–